

CS 581 – ADVANCED ARTIFICIAL INTELLIGENCE

TOPIC: CLASSIFICATION



Mustafa Bilgic



<http://www.cs.iit.edu/~mbilgic>



<https://twitter.com/bilgicm>

CLASSIFICATION

- AI → ML → Supervised Learning → Classification
- Given a dataset $D = \{\langle X_i, y_i \rangle\}$ where
 - X_i is the input
 - y_i is the discrete-valued output
- Learn a function $f(X_j) \rightarrow y_j$
- We would like f to **generalize** to **unseen** data
 - As opposed to memorizing the given data

CLASSIFICATION EXAMPLES

- Email classification
- Medical diagnosis
- Face recognition
- Optical character/digit recognition
- Sentiment classification
- ...

ALGORITHMS

- Decision trees
- Nearest neighbor classification
- Naïve Bayes
- Logistic regression
- Support vector machines
- Neural networks
- ...

GENERALIZATION

- The purpose of f
 - Is not to memorize the “seen” data
 - Is to generalize to “unseen” data
- We need a performance metric
- We need to test f ’s performance on a dataset that it has not seen

TYPES OF ERRORS – CLASSIFICATION

- Pick which cases should be called “positive”
 - Spam, HasHeartDisease, etc.
- *False positive*
 - Falsely classifying an object as positive
 - E.g., classifying a legitimate email as spam, diagnosing a healthy patient as having heart disease, etc.
 - Also called *Type I* error
- *False negative*
 - Falsely classifying an object as negative
 - E.g., classifying a spam email as not-spam, claiming that a heart-disease patient is healthy, etc.
 - Also called *Type II* error

A FEW PERFORMANCE MEASURES

- 0/1 loss; error or accuracy
- Precision
- Recall
- F1
- Log-loss
- Fairness
- Domain-specific performance measures
- ...

CONFUSION MATRIX

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

ACCURACY

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

$$Accuracy = \frac{Num\ Correct}{Data\ Size} = \frac{TP + TN}{TP + TN + FP + FN}$$

PRECISION

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

$$Precision = \frac{True\ Positive}{Predicted\ Positive} = \frac{TP}{TP + FP}$$

TRUE POSITIVE RATE – RECALL – SENSITIVITY

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

$$TPR = Recall = \frac{\text{True Positive}}{\text{Actual Positive}} = \frac{TP}{TP + FN}$$

F1

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

$$F1 = \frac{2 * Precision * Recall}{Precision + Recall}$$

$$\frac{\frac{1}{P} + \frac{1}{R}}{2}$$

TRUE NEGATIVE RATE – SPECIFICITY

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

$$TNR = Specificity = \frac{True\ Negative}{Actual\ Negative} = \frac{TN}{TN + FP}$$

FALSE POSITIVE RATE – FALL-OUT

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

$$FPR = FallOut = \frac{False\ Positive}{Actual\ Negative} = \frac{FP}{TN + FP}$$

FALSE NEGATIVE RATE – MISS RATE

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

$$FNR = Miss Rate = \frac{False\ Negative}{Actual\ Positive} = \frac{FN}{TP + FN}$$

SPLITTING THE DATASET

1. Train-test splits
2. Train-validation-test splits
3. Cross-validation

TRAIN-TEST SPLIT

- Randomly split the data into two disjoint sets
- A typical approach: $2/3$ for train and $1/3$ for test
- Train your model on training data and evaluate it on the test data
 - Use your favorite performance metric
- Report your performance as the expected performance on unseen data
- Caveats
 - You need a large dataset for this to work
 - You cannot tune your parameters on the test data

TRAIN-VALIDATION-TEST SPLIT

- Split your data into three disjoint sets
 - Train, validation, test
- Train your model(s) on the training data
- Evaluate your model(s) on the validation data
- Pick the model that performs best on the validation data
- Test the model on the test data, and report its performance
- Caveat
 - You need a really big dataset for this to work

CROSS-VALIDATION

- Split your data into k disjoint sets
- Each time, one set is the test set and the rest is the training set
- See OneNote for more detailed explanation and illustration

REAL LIFE MEASURES

- Not as clean as the ones we discussed
- Consider self-driving cars, medical diagnosis, crime prediction, fraud detection, and so on
- Often, there is not a single performance measure
- Performance is handled on a case-by-case basis; not on an aggregate level