

# CS 581 – ADVANCED ARTIFICIAL INTELLIGENCE

TOPIC: CLASSIFICATION



**Mustafa Bilgic**



<http://www.cs.iit.edu/~mbilgic>



<https://twitter.com/bilgicm>

# CLASSIFICATION

- AI → ML → Supervised Learning → Classification
- Given a dataset  $D = \{\langle X_i, y_i \rangle\}$  where
  - $X_i$  is the input
  - $y_i$  is the discrete-valued output
- Learn a function  $f(X_j) \rightarrow y_j$
- We would like  $f$  to **generalize** to **unseen** data
  - As opposed to memorizing the given data

# CLASSIFICATION EXAMPLES

- Email classification
- Medical diagnosis
- Face recognition
- Optical character/digit recognition
- Sentiment classification
- ...

# ALGORITHMS

- Decision trees
- Nearest neighbor classification
- Naïve Bayes
- Logistic regression
- Support vector machines
- Neural networks
- ...

# GENERALIZATION

- The purpose of  $f$ 
  - Is not to memorize the “seen” data
  - Is to generalize to “unseen” data
- We need a performance metric
- We need to test  $f$ ’s performance on a dataset that it has not seen

# TYPES OF ERRORS – CLASSIFICATION

- Pick which cases should be called “positive”
  - Spam, HasHeartDisease, etc.
- *False positive*
  - Falsely classifying an object as positive
    - E.g., classifying a legitimate email as spam, diagnosing a healthy patient as having heart disease, etc.
  - Also called *Type I* error
- *False negative*
  - Falsely classifying an object as negative
    - E.g., classifying a spam email as not-spam, claiming that a heart-disease patient is healthy, etc.
  - Also called *Type II* error

# A FEW PERFORMANCE MEASURES

- 0/1 loss; error or accuracy
- Precision
- Recall
- F1
- Log-loss
- Regret
- Fairness
- Domain-specific performance measures
- ...

# CONFUSION MATRIX

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative



# ACCURACY

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

$$Accuracy = \frac{Num\ Correct}{Data\ Size} = \frac{TP + TN}{TP + TN + FP + FN}$$

# PRECISION

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

$$Precision = \frac{True\ Positive}{Predicted\ Positive} = \frac{TP}{TP + FP}$$

## TRUE POSITIVE RATE – RECALL – SENSITIVITY

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

$$TPR = Recall = \frac{\text{True Positive}}{\text{Actual Positive}} = \frac{TP}{TP + FN}$$

# F1

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

$$F1 = \frac{2 * Precision * Recall}{Precision + Recall}$$

## TRUE NEGATIVE RATE – SPECIFICITY

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

$$TNR = Specificity = \frac{True\ Negative}{Actual\ Negative} = \frac{TN}{TN + FP}$$

## FALSE POSITIVE RATE – FALL-OUT

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

$$FPR = FallOut = \frac{False\ Positive}{Actual\ Negative} = \frac{FP}{TN + FP}$$

## FALSE NEGATIVE RATE – MISS RATE

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

$$FNR = Miss Rate = \frac{False\ Negative}{Actual\ Positive} = \frac{FN}{TP + FN}$$

# SPLITTING THE DATASET

1. Train-test splits
2. Train-validation-test splits
3. Cross-validation



# TRAIN-TEST SPLIT

- Randomly split the data into two disjoint sets
- A typical approach:  $2/3$  for train and  $1/3$  for test
- Train your model on training data and evaluate it on the test data
  - Use your favorite performance metric
- Report your performance as the expected performance on unseen data
- Caveats
  - You need a large dataset for this to work
  - You cannot tune your parameters on the test data

# TRAIN-VALIDATION-TEST SPLIT

- Split your data into three disjoint sets
  - Train, validation, test
- Train your model(s) on the training data
- Evaluate your model(s) on the validation data
- Pick the model that performs best on the validation data
- Test the model on the test data, and report its performance
- Caveat
  - You need a really big dataset for this to work

# CROSS-VALIDATION

- Split your data into  $k$  disjoint sets
- Each time, one set is the test set and the rest is the training set
- See OneNote for more detailed explanation and illustration

# REAL LIFE MEASURES

- Not as clean as the ones we discussed
- Consider self-driving cars, medical diagnosis, crime prediction, fraud detection, and so on
- Often, there is not a single performance measure
- Performance is handled on a case-by-case basis; not on an aggregate level