

# CS 581 – ADVANCED ARTIFICIAL INTELLIGENCE

TOPIC: DECISION MAKING UNDER UNCERTAINTY



**Mustafa Bilgic**



<http://www.cs.iit.edu/~mbilgic>



<https://twitter.com/bilgicm>

# MOTIVATION

- The agent needs reason in an uncertain world
- Uncertainty can be due to
  - Noisy sensors (e.g., temperature, GPS, camera, etc.)
  - Imperfect data (e.g., low resolution image)
  - Missing data (e.g., lab tests)
  - Imperfect knowledge (e.g., medical diagnosis)
  - Exceptions (e.g., all birds fly except ostriches, penguins, birds with injured wings, dead birds, ...)
  - Changing data (e.g., flu seasons, traffic conditions during rush hour, etc.)
  - ...
- The agent still must act (e.g., step on the breaks, diagnose a patient, order a lab test, ...)

# WHAT TO DO?

- Probability of a given email being spam is 0.6. Where should that email be delivered?
- Probability of rain is 0.6. I don't like carrying my umbrella around, but I also don't like getting wet. Should I take my umbrella?
- Probability of a given patient suffering from a heart disease is 0.6. Should more tests be conducted and if so, which ones and in which order? What should the treatment plan be?
- Probability of the product selling is 0.6. If we spend \$100K in ads, the probability goes up to 0.7. Is \$100K on ads justified?

# RATIONAL AGENT

- Given world states, a utility function, actions, transitions, evidence, and probabilities,
- A rational agent chooses the action that maximizes expected utility

$$action = \operatorname{argmax}_a EU(a|e)$$

# UNCERTAINTY REPRESENTATION

## ○ Covered

- Probability background (marginal, conditional, independence, Bayes rule, ...)
- Probabilistic classification (naïve Bayes, logistic regression, neural networks)

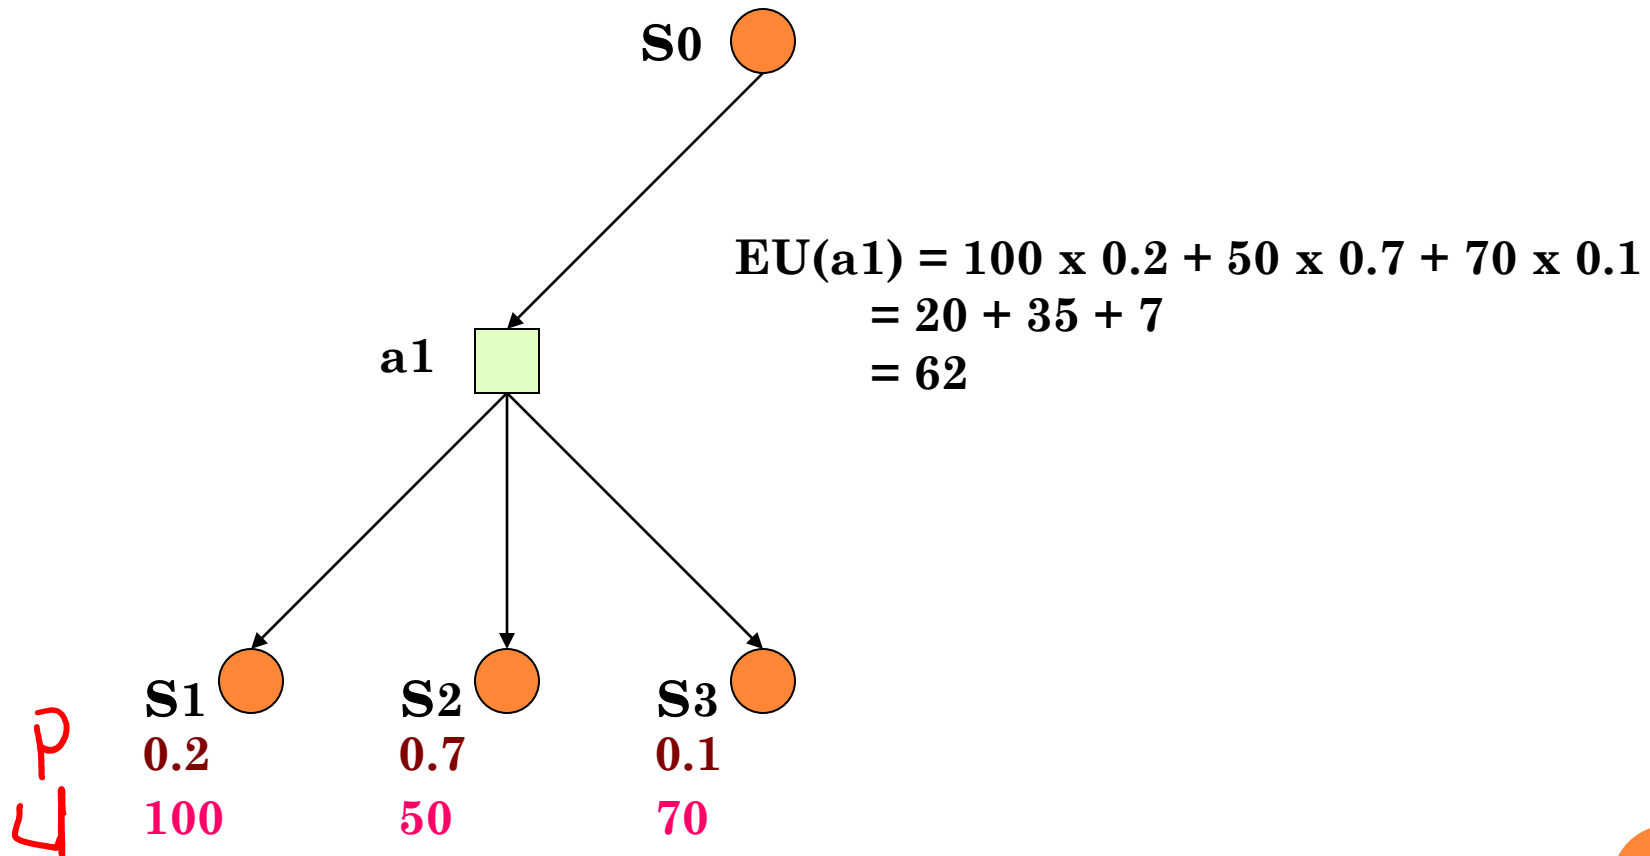
## ○ Skipped

- Bayesian networks (covered in CS 480 and CS 583)

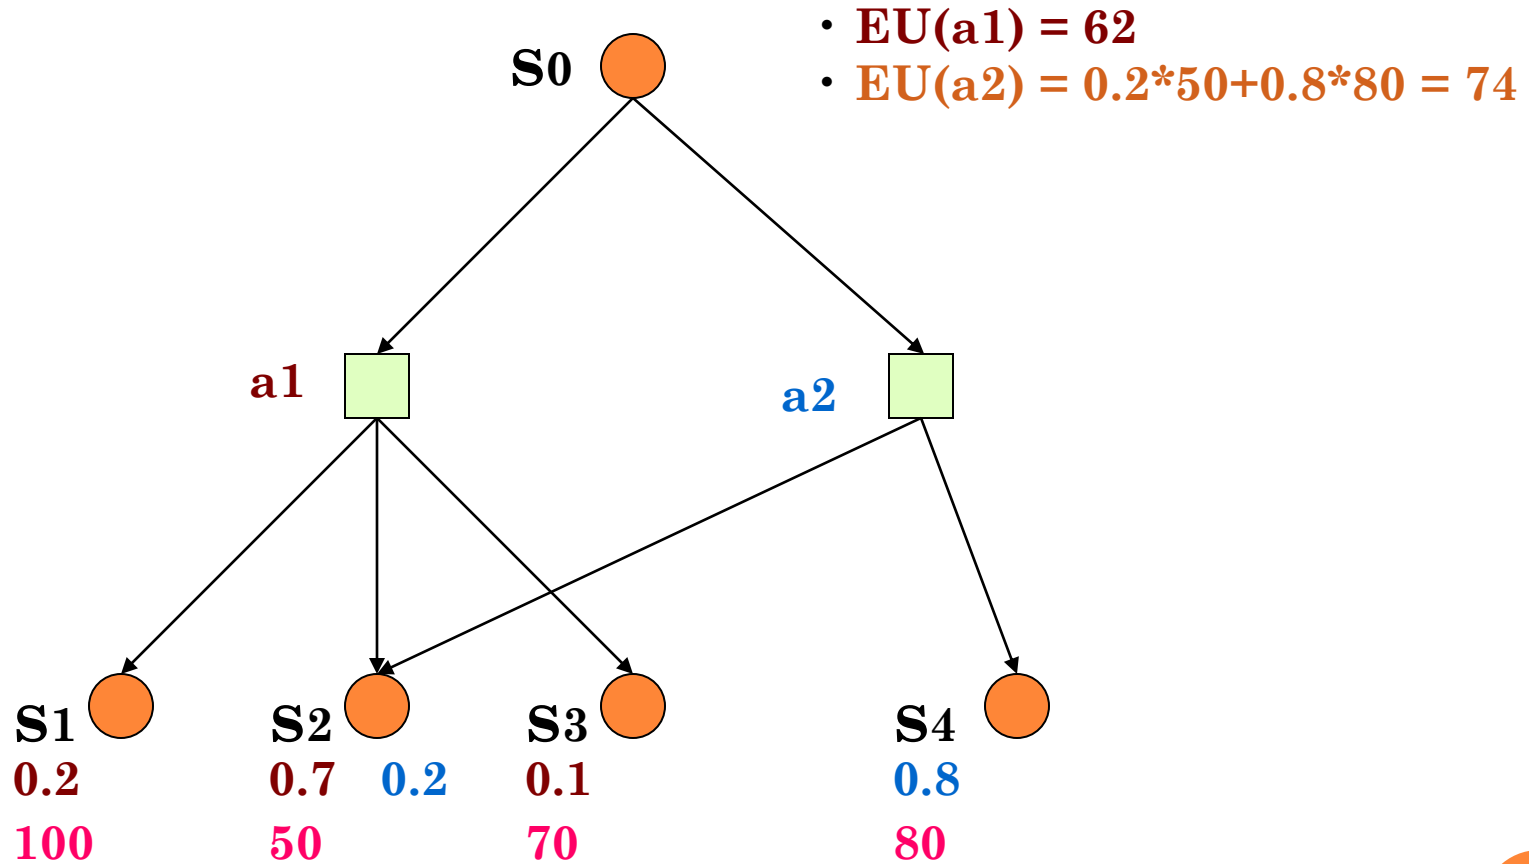
# UTILITY THEORY

- Lottery:  $n$  possible outcomes with probabilities
  - $[p_1, S_1; p_2, S_2; \dots p_n, S_n]$
  - Each  $S_i$  can be an atomic state or another lottery
- **Expected utility of a lottery**
  - $EU([p_1, S_1; p_2, S_2; \dots p_n, S_n]) = \sum_{i=1}^n p_i U(S_i)$

# ONE ACTION EXAMPLE



# TWO ACTIONS EXAMPLE





# UTILITY $\neq$ MONEY

- Most agents prefer more money to less money,
  - But this does not mean money behaves as a utility function
- For example, which lottery would you prefer
  - $L_1$ : [1, \$1 Million]  $\sim 1M$
  - $L_2$ : [0.5, \$0; 0.5, \$2.5 Million]  $1.25M$
- If money served as a utility function, then you'd prefer  $L_2$  no matter what, but the answer *often* depends on how much money you currently have
  - The utility of money depends on what you prefer
    - If you are short on cash, a little more *certain* money can help
    - If you are already billionaire, you might take the risk
    - Or if you are swimming in debt, you might like to gamble

# UTILITY $\neq$ MONEY

- Let's say you currently have \$k and let  $S_k$  represent the state of having \$k
- $EU(L_1) = U(S_{k+1M})$
- $EU(L_2) = 0.5 * U(S_k) + 0.5 * U(S_{k+2.5M})$
- The rational choice depends on your preferences for  $S_k$ ,  $S_{k+1M}$ , and  $S_{k+2.5M}$ 
  - i.e, it depends on the values of  $U(S_k)$ ,  $U(S_{k+1M})$ , and  $U(S_{k+2.5M})$
- $U(S_i)$  does not have to be a linear function of i, and for people it often is not

# WHICH ACTION TO TAKE?

## ○ Given

- A probability distribution
- Choice of actions and their effects on the distribution
- Evidence
- A utility function (as a function of states and possibly actions)

## ○ Find

- Which action maximizes the expected utility?

# A SIMPLE SPAM FILTERING EXAMPLE

- Given an email:
  - $P(s | e) = 0.6$
  - $P(\sim s | e) = 0.4$
- Actions:
  - deliverIntoSpamFolder (dS)
  - deliverIntoInboxFolder (dI)
- Utility function:
  - $U(dS, s) = 200$
  - $U(dI, \sim s) = 100$
  - $U(dI, s) = -100$
  - $U(dS, \sim s) = -500$
- Where should this email be delivered and why?

$$t \Rightarrow p \times c + (1-p) \times b$$

## UMBRELLA EXAMPLE

$$\neg t \Rightarrow p \times d + (1-p) \times a$$

- Check the weather:

- $P(r) = p$
- $P(\sim r) = 1-p$

- Actions:

- takeUmbrella (t)
- $\sim$ takeUmbrella ( $\sim t$ )

- Utility function:

- $U(\sim r, \sim t) = a$
- $U(\sim r, t) = b$
- $U(r, t) = c$
- $U(r, \sim t) = d$

- When should you take the umbrella?

# VALUE OF INFORMATION

- If I can buy more information to help with my decision, up to how much should I pay for that information?
  - Currency here is in “utility” units
- Value of information
  - Expected utility after the information is acquired
    - Minus
  - Expected utility before the information is acquired
- There is one catch: we do not know the content of the information before we acquire it
  - Solution: take an expectation over the possible outcomes

# VALUE OF INFORMATION – ALL BINARY

- Probability of state:
  - $P(s) = p, P(\sim s) = 1-p$
- Actions
  - $a, \sim a$
- Additional info,  $E$ 
  - $P(s | e) = q, P(\sim s | e) = 1-q$
  - $P(s | \sim e) = r, P(\sim s | \sim e) = 1-r$
  - $P(e) = v, P(\sim e) = 1-v$
- Utilities
  - $U(s, a) = u1$
  - $U(s, \sim a) = u2$
  - $U(\sim s, a) = u3$
  - $U(\sim s, \sim a) = u4$
- What is the value of  $E$ ?  $VOI(E)$ ?

# NEXT

- Multi-armed bandit
- Markov decision processes
- Reading
  - Chapter 17 of the AI book by Stuart & Russel (<http://aima.cs.berkeley.edu/>)
  - Chapters 2 and 3 of the RL book by Sutton & Barto (<http://www.incompleteideas.net/book/the-book-2nd.html>)



# MULTI-ARMED BANDIT

# SETUP

- K slot machines
  - Each one is a one-armed bandit
- Unknown reward functions
  - The distribution of rewards for each machine are unknown
- Limited resources
  - Have a limited number of tries (or alternatively, future rewards are discounted)
- Can gather information
  - Each try gives a (potentially) zero/negative reward, but also is useful for information gathering purposes
- Main objective
  - Maximize rewards
- Exploration vs exploitation trade-off
  - How should you balance exploitation (sticking with a machine that looks good) versus exploration (trying new machines)?

# NOTATION

- $A_t$ : action at time  $t$
- $R_t$ : reward at time  $t$
- Sequence of actions and rewards:  
 $A_1, R_1, A_2, R_2, \dots, A_T, R_T$
- $q_*(a) = E[R_t \mid A_t = a]$ 
  - Not given; otherwise, the solution is trivial
- $Q_t(a)$ : Average reward for action  $a$ , up to, excluding, time  $t$ 
  - Estimated using prior experiences
  - If action  $a$  has never been taken before time  $t$ ,  $Q_t(a)$  is initialized to a default value

# OPTIMAL STRATEGY

- $\operatorname{argmax}_a q_*(a)$
- At time  $t$ , choose the action that has the highest expected reward
- Challenge
  - We do not know  $q_*(a)$

# GREEDY STRATEGY (EXPLOIT ONLY)

- $\operatorname{argmax}_a Q_t(a)$
- Calculate the average reward for each action, up to time  $t$ , and choose the maximum one
- Problems
  - Initial values of  $Q_t(a)$  plays a big role
  - It simply tries the best action it has found; purely exploitation
  - Could easily miss other actions that are better but not tried


# RANDOM STRATEGY (EXPLORE ONLY)

- Choose an action at random
- Good
  - Explores all the time
- Bad
  - Does not exploit; does not learn

# EPSILON GREEDY

- Given an exploration parameter  $\epsilon$
- At each step  $t$ :
  - With probability  $\epsilon$ , choose a random action (explore)
  - With probability  $1 - \epsilon$ , choose the current best action:  
 $\operatorname{argmax}_a Q_t(a)$  (exploit)
- $\epsilon = 0$  is fully greedy, and  $\epsilon = 1$  is fully random

# SIMULATION SETUP

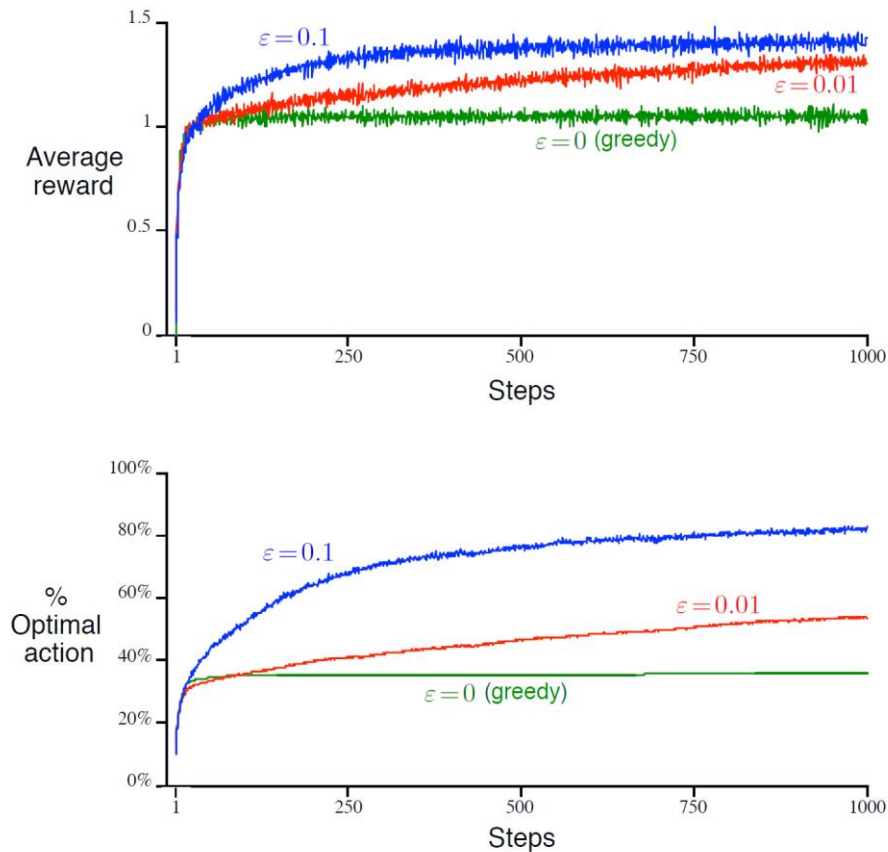


**Figure 2.1:** An example bandit problem from the 10-armed testbed. The true value  $q_*(a)$  of each of the ten actions was selected according to a normal distribution with mean zero and unit variance, and then the actual rewards were selected according to a mean  $q_*(a)$ , unit-variance normal distribution, as suggested by these gray distributions.

Figure from <http://www.incompleteideas.net/book/the-book-2nd.html>



# SIMULATION RESULTS




**Figure 2.2:** Average performance of  $\epsilon$ -greedy action-value methods on the 10-armed testbed. These data are averages over 2000 runs with different bandit problems. All methods used sample averages as their action-value estimates.

Figure from <http://www.incompleteideas.net/book/the-book-2nd.html>

# UCB

- The upper-confidence bound (UCB) method calculates the upper bound on the mean for each slot and chooses the machine with max value
  - $Q_t(a) + c \sqrt{\frac{\ln(t)}{N_t(a)}}$ , where  $N_t(a)$  is the number of times the action  $a$  is tried and  $c$  is the trade-off parameter
- The term in the square root is a measure of uncertainty in  $Q_t(a)$ ; and hence the name upper confidence bound
  - The upper confidence is derived using the Chernoff-Hoeffding bound
- The exploration grows with  $\ln(t)$ , shrinks with  $N_t(a)$
- We've seen this before; where?

# SIMULATION RESULTS



**Figure 2.4:** Average performance of UCB action selection on the 10-armed testbed. As shown, UCB generally performs better than  $\epsilon$ -greedy action selection, except in the first  $k$  steps, when it selects randomly among the as-yet-untried actions.

# A RUNNING AVERAGE COMPUTATION

$$\begin{aligned}Q_{n+1} &= \frac{1}{n} \sum_{i=1}^n R_i \\&= \frac{1}{n} \left( R_n + \sum_{i=1}^{n-1} R_i \right) \\&= \frac{1}{n} \left( R_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} R_i \right) \\&= \frac{1}{n} \left( R_n + (n-1) Q_n \right) \\&= \frac{1}{n} \left( R_n + n Q_n - Q_n \right) \\&= Q_n + \frac{1}{n} \left[ R_n - Q_n \right],\end{aligned}$$

Figure from <http://www.incompleteideas.net/book/the-book-2nd.html>

# UPDATE RULE

- $Q_{n+1} = Q_n + \frac{1}{n}(R_n - Q_n)$
- $\text{new} = \text{old} + \text{stepSize} * (\text{target} - \text{old})$
- For computing the exact average, stepSize is a function of  $n$

# NON-STATIONARY REWARDS

- What if the reward distribution changes over time?
- We'd like to give recent rewards more weight
- A simple approach
  - $Q_{n+1} = Q_n + \alpha(R_n - Q_n)$
- Earlier rewards have lower weights

- $$\begin{aligned} Q_{n+1} &= Q_n + \alpha(R_n - Q_n) \\ &= \alpha R_n + (1 - \alpha)Q_n \\ &= \alpha R_n + (1 - \alpha)[\alpha R_{n-1} + (1 - \alpha)Q_{n-1}] \end{aligned}$$

# MARKOV DECISION PROCESSES

# PROBLEM SETTING

- The world is represented through states
- At each state, an agent is given 0 (terminal states) or more actions to choose from
- Each action moves the agent, probabilistically, to a state (could be the current state) and results, probabilistically, in a reward (could be zero, negative, positive)
- The agent needs to maximize the sum of the rewards it accumulates over time
- Greedy strategy with respect to immediate rewards often do not work; the agent needs to consider the long-term consequences of its actions




# MARKOV DECISION PROCESS

- A sequential decision-making process
- Stochastic environment
- Markov transition model
- Additive rewards
- Applications: planning
  - Search Google
  - <http://www.it.uu.se/edu/course/homepage/aism/st11/MDPApplications1.pdf>

# MDP DIAGRAM EXAMPLE


- See OneNote



# NOTATION

- $P(s' | s, a)$  Probability of arriving at state  $s'$  given we are at state  $s$  and take action  $a$
- $R(s, a, s')$  The reward the agent receives when it transitions from state  $s$  to state  $s'$  via action  $a$
- $\pi(s)$  The action recommended by policy  $\pi$  at state  $s$
- $\pi^*$  Optimal policy
- $U^\pi(s)$  The expected utility obtained via executing policy  $\pi$  starting at state  $s$
- $U^{\pi^*}(s)$  is often abbreviated as  $U(s)$
- $Q^\pi(s, a)$  The expected utility of taking action  $a$  at state  $s$
- $\gamma$  Discount factor  $[0, 1]$

# RUNNING EXAMPLE



**Figure 17.1** (a) A simple, stochastic  $4 \times 3$  environment that presents the agent with a sequential decision problem. (b) Illustration of the transition model of the environment: the “intended” outcome occurs with probability 0.8, but with probability 0.2 the agent moves at right angles to the intended direction. A collision with a wall results in no movement. Transitions into the two terminal states have reward +1 and -1, respectively, and all other transitions have a reward of -0.04.

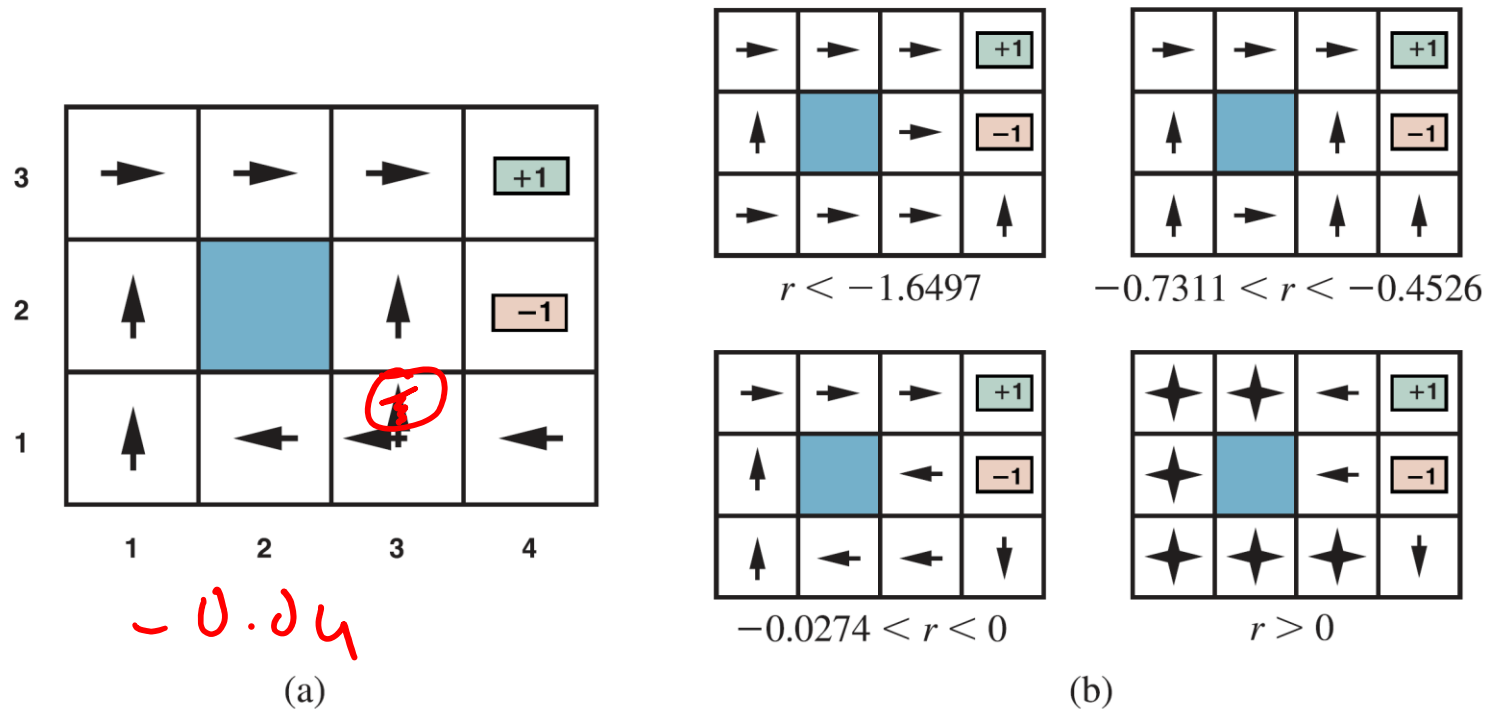
Figure from <http://aima.cs.berkeley.edu/figures.pdf>

# SOLUTION?

- A fixed action sequence is not the answer due to stochasticity
  - For example, [Up, Up, Right, Right, Right] is not a solution
  - It would be a solution if the environment was deterministic
- A solution must specify the agent should do in any state that the agent might reach
  - This is called a **policy**
- Policy notation:  $\pi$ 
  - $\pi(s)$  specifies what action the agent should take at state  $s$
- An **optimal policy** is the one that maximizes the expected utility
  - $\pi^*$

# RUNNING EXAMPLE

$$\gamma = 1$$



**Figure 17.2** (a) The optimal policies for the stochastic environment with  $r = -0.04$  for transitions between nonterminal states. There are two policies because in state (3,1) both *Left* and *Up* are optimal. (b) Optimal policies for four different ranges of  $r$ .

Figure from <http://aima.cs.berkeley.edu/figures.pdf>

# UTILITY OF STATES

- The agent receives a reward at each state
- Utility of a state  $s$  given a policy  $\pi$  is the expected reward that the agent will get starting from state  $s$  and taking actions according to policy  $\pi$
- Let  $S_t$  denote the state that the agent reaches at time  $t$
- $U^\pi(s) = E[\sum_{t=0}^{\infty} \gamma^t R(S_t, \pi(S_t), S_{t+1})]$
- The expectation is with respect to the transition probabilities

## U(s) VS R(s, A, s')

- $R(s, a, s')$  is the short-term immediate reward the agent receives when it transitions from state  $s$  to state  $s'$  via action  $a$
- $U(s)$  is the long-term cumulative reward from  $s$  onward
- $U^\pi(s) = E[\sum_{t=0}^{\infty} \gamma^t R(S_t, \pi(S_t), S_{t+1})]$



# BELLMAN EQUATION

$$U^{\pi}(s) = \sum_{s'} P(s'|s, \pi(s)) \times [R(s, \pi(s), s') + \gamma U^{\pi}(s')]$$

# BELLMAN OPTIMALITY EQUATION

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U(s')]$$

# THE OPTIMAL POLICY

- The optimal policy is the one that maximizes the expected utility
  - $\pi_s^* = \operatorname{argmax}_{\pi} U^{\pi}(s)$
- Remember that  $\pi_s^*$  is a policy; that is, it recommends an action for each state, regardless of whether it is the starting state or not
- It is optimal when the starting state is  $s$
- When the rewards are discounted, the optimal policy is independent of the start state
  - The optimality of the policy does not depend on the starting state but of course the action sequence depends on the starting state
- True utility of each state is defined as  $U^{\pi^*}(s)$  -- the expected rewards the agent will receive if it executes the optimal policy starting at  $s$

## ACTION-UTILITY FUNCTION $Q(s, a)$

- $Q^\pi(s, a)$  The expected utility of taking action  $a$  at state  $s$  and then following policy  $\pi$
- $Q^\pi(s, a) = \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U^\pi(s')]$
- $U^{\pi^*}(s) = \max_{a \in A(s)} Q^{\pi^*}(s, a)$
- $Q^{\pi^*}(s, a) = \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma \max_{a' \in A(s')} Q^{\pi^*}(s', a')]$ 
  - Bellman optimality equation for the Q function

$$\pi^*(s) = \operatorname{argmax}_{a \in A(s)} Q^{\pi^*}(s, a)$$

# RUNNING EXAMPLE

$$V(T_s) = 0$$

---

3	0.8516	0.9078	0.9578	<del>+1</del> <b>A</b>
2	0.8016		0.7003	<del>-1</del> <b>B</b>
1	0.7453	0.6953	0.6514	0.4279
	1	2	3	4

**Figure 17.3** The utilities of the states in the  $4 \times 3$  world with  $\gamma = 1$  and  $r = -0.04$  for transitions to nonterminal states.

---

Figure from <http://aima.cs.berkeley.edu/figures.pdf>

# Q(S, A)

Highlighted values are inaccurate.

					UP			
0.8516	0.9078	0.9578	1		0.81722	0.86718	0.91702	0
0.8016		0.7003	-1		0.8016		0.70027	0
0.7453	0.6953	0.6514	0.4279		0.74534	0.65591	0.63256	-0.70007
					DOWN			
					0.77722	0.86718	0.71102	0
					0.71656		0.45515	0
					0.7003	0.65591	0.59344	0.41025
					LEFT			
					0.8066	0.82284	0.85205	0
					0.76097		0.68116	0
					0.71093	0.6953	0.65141	0.42791
					RIGHT			
					0.85156	0.9078	0.95781	0
					0.76097		-0.64708	0
					0.67093	0.62018	0.43749	0.24911
					MAX			
					0.85156	0.9078	0.95781	0
					0.8016		0.70027	0
					0.74534	0.6953	0.65141	0.42791

# EXERCISE

- Confirm that the utilities given in the previous slide satisfy the Bellman equations

# HOW TO FIND $\pi^*$

- $\pi^*(s) = \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U^{\pi^*}(s')]$
- However, we are not given  $U^{\pi^*}(s)$
- Two algorithms for finding optimal policies
  1. Value iteration
  2. Policy iteration



# VALUE ITERATION

- $U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U(s')]$   
*(Handwritten red annotations: a red 'i+1' under the first 'i' in 'U(s)' and a red '(' under the 'U' in 'U(s')')*
- $n$  possible states,  $n$  Bellman equations, one for each state
- However, these are non-linear equations, due to the max operator
- One approach: iterative
  - Start with an initial guess (could be random)
  - Iterate until convergence

# POLICY ITERATION

- Start with an initial policy  $\pi_0$
- Alternate between
  1. Policy evaluation: given policy  $\pi_i$ , calculate  $U^{\pi_i}$ 
    - Can be calculated exactly (linear equations) or iteratively
  2. Policy improvement: Calculate a new MEU policy  $\pi_{i+1}$ , using the utilities calculated in the previous step
- Stop when policy no longer changes

# FINDING OPTIMAL POLICIES

- The previous two slides computed U
  - Many people find the concept of U more intuitive than the concept of Q
- Using Q to find the optimal policy makes more sense
  - $\pi^*(s) = \operatorname{argmax}_a Q^{\pi^*}(s, a)$

# NEXT

## ○ Reinforcement learning

- In fact, we already covered many of the fundamentals of RL
  - Value iteration, policy iteration, exploration vs exploitation trade-off
- We are now ready to make the leap from MDPs to RL
- RL can be considered as solving an MDP where the transition and reward dynamics are unknown